

Introduction

I'm a massive fan of the Python language. It was the first programming language I learned for data science and it has been a constant companion ever since. Three things stand out about Python for me:

- Its ease and flexibility
- Industry-wide acceptance: It is far and away the most popular language for data science in the industry
- The sheer number of Python libraries for data science

In fact, there are so many Python libraries out there that it can become overwhelming to keep abreast of what's out there. That's why I decided to take away that pain and compile this list of 24 awesome Python libraries covering the end-to-end data science lifecycle.

Python libraries for different data science tasks:

- Python Libraries for Data Collection
 - Beautiful Soup
 - Scrapy
 - Selenium
- Python Libraries for Data Cleaning and Manipulation
 - Pandas
 - PyOD
 - NumPy
 - Spacy
- Python Libraries for Data Visualization
 - Matplotlib
 - Seaborn
 - Bokeh
- Python Libraries for Modeling
 - Scikit-learn
 - TensorFlow
 - PyTorch
- Python Libraries for Model Interpretability
 - Lime
 - H2O
- Python Libraries for Audio Processing
 - Librosa
 - Madmom
 - pyAudioAnalysis
- Python Libraries for Image Processing
 - OpenCV-Python
 - Scikit-image
 - Pillow
- Python Libraries for Database
 - Psycopg
 - SQLAlchemy
- Python Libraries for Deployment
 - Flask